

# LINEAR REGRESSION

```

In [3]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

data={
    "pressure": [0.0002, 0.0012, 0.0060, 0.0300, 0.0900, 0.1000, 0.1100, 0.1200, 0.2
700, 36.0000, 47.0000, 55.0000, 73.0000, 75.0000, 80.7000, 120.7000, 172.0000, 220.0
000, 275.0000, 320.0000],
    "temperature": [0, 20, 40, 60, 80, 100, 120, 140, 160, 180, 200, 220, 240, 260, 280, 30
0, 320, 322, 340, 342]
}

df=pd.DataFrame(data)

x=df.iloc[:,0].values.reshape(-1,1)
y=df.iloc[:,1].values.reshape(-1,1)

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_sta
te=42)

model=LinearRegression()
model.fit(x_train,y_train)
print(model.predict([[100]]))
y_pred = model.predict(x_test)
mae=mean_absolute_error(y_test,y_pred)
mse=mean_squared_error(y_test,y_pred)
r2=r2_score(y_test,y_pred)
print("mse: ",mse)
print("mae: ",mae)
print("R2: ",r2)


x_line = np.linspace(x.min(), x.max(), 100).reshape(-1, 1)
plt.plot(x_line, model.predict(x_line), 'r')
plt.scatter(x,y)
plt.show()

```

```

[[217.09926859]]
mse: 8689.589903453918
mae: 78.86069844924648
R2: 0.6182204055905927

```



```
In [22]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Create custom dataset
pressure = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
temperature = [15, 18, 21, 23, 25, 28, 30, 33, 35, 38]

df = pd.DataFrame({
    "pressure": pressure,
    "temperature": temperature
})

# Prepare the data
X = df[["pressure"]].values # 2D array of features
y = df["temperature"].values.reshape(-1, 1) # target values

# Split into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

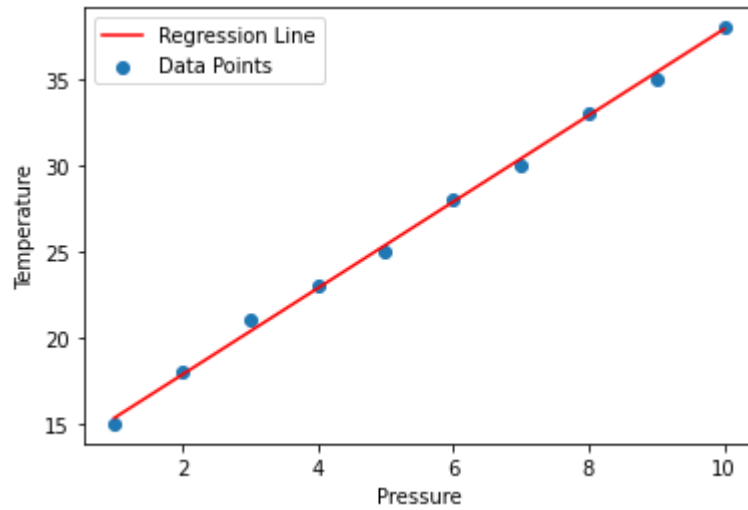
# Make predictions and evaluate
y_pred = model.predict(X_test)
print("MSE:", mean_squared_error(y_test, y_pred))
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R2:", r2_score(y_test, y_pred))

# Plot the regression line and data points
x_line = np.linspace(min(pressure), max(pressure), 100).reshape(-1, 1)
plt.plot(x_line, model.predict(x_line), 'r', label="Regression Line")
plt.scatter(X, y, label="Data Points")
plt.xlabel("Pressure")
plt.ylabel("Temperature")
plt.legend()
plt.show()
```

MSE: 0.09412158145065727

MAE: 0.28017241379310853

R2: 0.9986972791494718



## LOGISTIC

```
In [4]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

df=sns.load_dataset("iris")
x=df[["sepal_length","sepal_width","petal_length","petal_width"]]
y=df[["species"]]

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
model=LogisticRegression(C=0.01)
model.fit(x_train,y_train)
pred=model.predict(x_test)

print("accuracy: ",accuracy_score(y_test,pred))
print("\nClassification_report: ",classification_report(y_test,pred))
cm=confusion_matrix(y_test,pred)

sns.heatmap(cm,annot=True,cmap="Blues")
plt.xlabel("predicted labels")
plt.ylabel("actual labels")
plt.show()
```


C:\CYrus\anaconda\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
return f(*args, **kwargs)
```

accuracy: 0.9333333333333333

Classification\_report:                      precision      recall    f1-score    support

setosa	1.00	1.00	1.00	10
versicolor	1.00	0.78	0.88	9
virginica	0.85	1.00	0.92	11
accuracy			0.93	30
macro avg	0.95	0.93	0.93	30
weighted avg	0.94	0.93	0.93	30



## DECISION TREE

```

In [5]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree

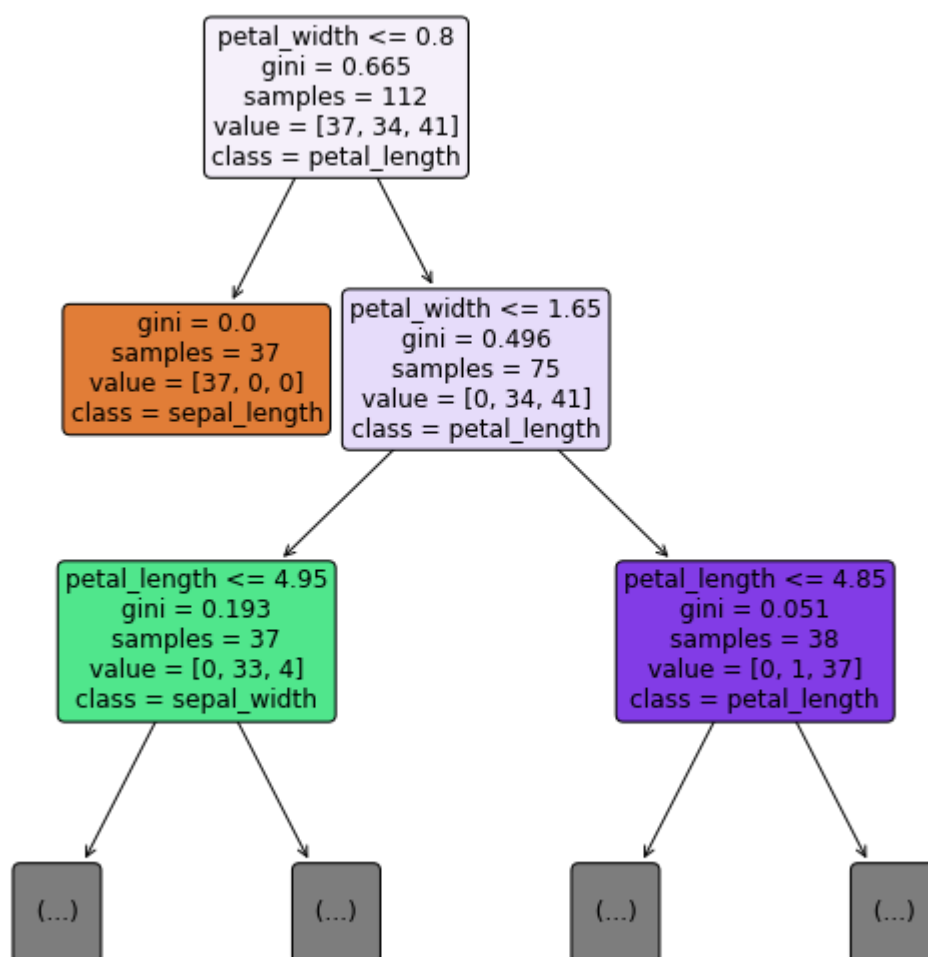
data=sns.load_dataset('iris')
x=data.values[:,0:4]
y=data.values[:,4]

x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=1)
model=DecisionTreeClassifier(random_state=54)
model.fit(x_train,y_train)

pred=model.predict(x_test)

plt.figure(figsize=(10,10))
plot_tree(model,filled=True,rounded=True,feature_names=data.columns[:4].to_
list(),class_names=data.columns[:5].unique().to_list(),max_depth=2)
plt.show()

```




## KMEANS CLUSTERING

```
In [6]: import seaborn as sns
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

data=sns.load_dataset('iris')
model=KMeans(n_clusters=3)
model.fit(data.drop('species',axis=1))
labels=model.labels_

plt.scatter(data['sepal_length'],data['sepal_width'],c=labels)
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')
plt.title('Iris Dataset Clustering')
plt.show()
```



## PERCEPTRON

```
In [8]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score
data = load_iris()
x, y = data.data[:100, :], data.target[:100]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
perceptron = Perceptron(random_state=42)
perceptron.fit(x_train, y_train)
y_pred = perceptron.predict(x_test)
print(f'Prediction: {y_pred}')
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

Prediction: [1 1 1 0 0 0 0 1 0 0 0 0 1 0 1 0 1 1 0 0]  
Accuracy: 1.0

## RANDOM FOREST CLASSIFIER



```

In [9]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt
iris = load_iris()
x = iris.data
y = iris.target
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
model = RandomForestClassifier()
clf=model.fit(x_train, y_train)
y_pred = clf.predict(x_test)
# print(y_pred)
accuracy=accuracy_score(y_test,y_pred)
print(f'Accuracy:{accuracy:.2f}')
print("Report: ", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, cmap="Blues")
plt.xlabel("Predicted labels")
plt.ylabel("True labels")
plt.title("Confusion matrix")
plt.show()


```

Accuracy:1.00

Report:                      precision      recall    f1-score      support

0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11

accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30



## VOTING CLASSIFIER

```
In [10]: from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

data = load_iris()
X = data.data
y = data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
log_clf = LogisticRegression(max_iter=200)
tree_clf = DecisionTreeClassifier(random_state=42)
knn_clf = KNeighborsClassifier()

# Hard Voting Classifier (majority vote)
hard_voting_clf = VotingClassifier(estimators=[
    ('log_clf', log_clf),
    ('tree_clf', tree_clf),
    ('knn_clf', knn_clf)
], voting='hard')

# Soft Voting Classifier (average probabilities)
soft_voting_clf = VotingClassifier(estimators=[
    ('log_clf', log_clf),
    ('tree_clf', tree_clf),
    ('knn_clf', knn_clf)
], voting='soft')

# Train and evaluate Hard Voting Classifier
hard_voting_clf.fit(X_train, y_train)
y_pred_hard = hard_voting_clf.predict(X_test)
hard_acc = accuracy_score(y_test, y_pred_hard)
print(f"Hard Voting Classifier Accuracy: {hard_acc:.4f}")

# Train and evaluate Soft Voting Classifier
soft_voting_clf.fit(X_train, y_train)
y_pred_soft = soft_voting_clf.predict(X_test)
soft_acc = accuracy_score(y_test, y_pred_soft)
print(f"Soft Voting Classifier Accuracy: {soft_acc:.4f}")
```

Hard Voting Classifier Accuracy: 1.0000

Soft Voting Classifier Accuracy: 1.0000

## APRIORI

```
In [11]: import pandas as pd
from mlxtend.frequent_patterns import apriori, association_rules
from mlxtend.preprocessing import TransactionEncoder

dataset = [
    ['milk', 'bread', 'butter'],
    ['bread', 'butter'],
    ['milk', 'bread'],
    ['milk', 'bread', 'butter'],
    ['bread', 'butter']
]
te = TransactionEncoder()
te_ary = te.fit(dataset).transform(dataset)
df = pd.DataFrame(te_ary, columns=te.columns_)
frequent_itemsets = apriori(df, min_support=0.6, use_colnames=True)
print("\nFrequent Itemsets:")
print(frequent_itemsets)
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.0)
print("\nAssociation Rules:")
print(rules)
print("\nNumber of association rules: ",len(rules))
```

Frequent Itemsets:

	support	itemsets
0	1.0	(bread)
1	0.8	(butter)
2	0.6	(milk)
3	0.8	(butter, bread)
4	0.6	(milk, bread)

Association Rules:

	antecedents	consequents	antecedent support	consequent support	support
0	(butter)	(bread)	0.8	1.0	0.8
1	(bread)	(butter)	1.0	0.8	0.8
2	(milk)	(bread)	0.6	1.0	0.6
3	(bread)	(milk)	1.0	0.6	0.6

	confidence	lift	leverage	conviction
0	1.0	1.0	0.0	inf
1	0.8	1.0	0.0	1.0
2	1.0	1.0	0.0	inf
3	0.6	1.0	0.0	1.0

Number of association rules: 4

## naive bayes

```
In [12]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

# Load the iris dataset
iris = load_iris()
X, y = iris.data, iris.target

# Split into training and test sets (70% train, 30% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initialize and train the Gaussian Naïve Bayes classifier
nb = GaussianNB()
nb.fit(X_train, y_train)

# Make predictions on the test set
y_pred = nb.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9777777777777777

## svm

```

In [20]: import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
    classification_report
)

iris = load_iris()
X, y = iris.data, iris.target

X_2d = X[:, :2]

X_train, X_test, y_train, y_test = train_test_split(
    X_2d, y, test_size=0.3, random_state=42
)
clf = svm.SVC(kernel='linear', random_state=42)
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

acc = accuracy_score(y_test, y_pred)
print(f"Accuracy: {acc:.2f}")

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

report = classification_report(y_test, y_pred)
print("Classification Report:")
print(report)

```

Accuracy: 0.80

Confusion Matrix:

```

[[19  0  0]
 [ 0  7  6]
 [ 0  3 10]]

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	0.70	0.54	0.61	13
2	0.62	0.77	0.69	13
accuracy			0.80	45
macro avg	0.78	0.77	0.77	45
weighted avg	0.81	0.80	0.80	45